

Structure-Based 3D Drug Discovery Against the SARS-CoV-2 Main Protease

*Validated Redocking, a Pose-Graph Neural Rescoring Benchmark, and the De Novo
Candidate MPRO-D1*

MEGA-PROGRAM-27, Item 12

Computational Biology and Biophysics Research Series

September 2026 - Version 2 (expanded edition with the GNN benchmark and MPRO-D1)

Executive Summary

WHAT WAS BUILT: an end-to-end, redock-validated structure-based discovery pipeline for SARS-CoV-2 Mpro - real AutoDock Vina 1.2.7, real PDB/PubChem data, statistical enrichment analysis, two neural rescorers, and a de novo candidate generator - shipped as tested code with this paper.

WHAT WAS FOUND: (1) the protocol is validated - redocking JUN8-76-3A into 7KX5 reproduces the crystal pose at 0.97 Å heavy-atom RMSD (gate 2.0 Å), affinity -9.32 kcal/mol. (2) A benchmark head-to-head: a pose-graph neural rescorer trained on re-docked poses lifts leave-one-out AUROC to 0.871 $\pm$ 0.050 (5 seeds) versus 0.679 for a descriptor MLP and 0.339 for the raw engine, on the same 15 labeled compounds. That 15-compound ladder is the pilot; the primary benchmark is the 59-compound scaffold-grouped bigscreen, where raw Vina sits at 0.534 and both learned rescorers reach 0.830 (MLP) / 0.833 $\pm$ 0.010 (2D-GNN) - learned rescoring, not a winner between learned families, is the verdict at this n. (3) A de novo candidate, MPRO-D1, docking at -8.76 kcal/mol under the validated protocol, maximally dissimilar to every screened drug (Tanimoto 0.16), verified absent from PubChem as an exact structure, with close analogs disclosed. (4) The baseline that anchors the benchmark, reported as measured: raw Vina affinity does NOT rank actives above inactives on this set (AUROC 0.339, below random; top-8 hypergeometric $p = 0.900$) - without that negative the rescorer numbers would be unanchored.

WHAT WOULD FALSIFY IT: the redock gate is binary and re-runnable; the AUROC numbers recompute from the shipped poses; MPRO-D1's prediction is decided by a crystal or ITC experiment. Small-n caveats are stated wherever n is small.

UPDATE - MERGED ARCHITECTURE BENCHMARK (results/ablation.json): every scorer retrained on the same 59-compound label set under the same five scaffold-grouped splits, with preregistered controls. The pose-graph GNN (0.782) does NOT beat the 2D bond-graph GNN (0.902) or its own coordinate controls (random 0.831, shuffled 0.830); the pose-geometry superiority claim from the $n = 15$ pilot is withdrawn. No Vina leakage (descriptor drop without the engine feature: 0.006 against a >0.05 gate). Random-label control sane

(0.525). The standing result: learned rescoring beats raw Vina (0.540) by roughly 0.4 AUROC whichever learned family is used; fusion (0.931) is numerically best, descriptively.

Abstract

This study delivers a validated structure-based discovery pipeline for the SARS-CoV-2 main protease (Mpro), a quantified head-to-head benchmark of learned rescoring against the empirical docking engine, and MPRO-D1, a novel de novo candidate with a live-verified novelty record. Using the real AutoDock Vina 1.2 engine, we validate the protocol by redocking the non-covalent inhibitor JUN8-76-3A into Mpro (PDB 7KX5): the docked pose reproduces the experimental binding mode at 0.97 Å heavy-atom RMSD, inside the community-standard 2.0 Å gate, at -9.32 kcal/mol. Docking a curated set of 22 real approved and clinical compounds with literature Mpro labels sets the baseline the benchmark must beat, and it is reported as measured: raw Vina affinity separates actives from inactives at AUROC 0.339 - below chance - and top-8 enrichment is absent (hypergeometric $p = 0.900$). A descriptor MLP rescorer recovers to 0.679 leave-one-out AUROC, and a pose-graph neural rescorer - message passing over the docked protein-ligand contact graph - reaches 0.871 $\pm$ 0.050 (five seeds) on the same labels, a quantified head-to-head benchmark of learned rescoring against the physics-empirical engine. Finally, a constrained de novo enumerator proposes MPRO-D1, a biphenyl-amide ligand docking at -8.76 kcal/mol, Tanimoto 0.16 against the nearest screened drug, verified as an exact-structure novelty against PubChem with its close analogs disclosed. Every structure, ligand and label is real; every figure regenerates from shipped code. The primary benchmark is the scaffold-grouped 59-compound label set; the original 15-compound ladder is retained as a pilot calibration, not the headline. A merged architecture benchmark on that label set (all scorers under one scaffold-split protocol with preregistered controls) overturns the pilot's pose-geometry edge: the pose-graph GNN (0.782) trails the 2D graph (0.902) and its own shuffled-coordinate controls, so the pose-geometry superiority claim is withdrawn and learned rescoring as such - not a winning architecture - is the supported result. What this is not: not a validated inhibitor, not a clinical candidate, and not a synthesis plan - MPRO-D1 is a docking-ranked hypothesis with a live-verified novelty record. The limitations - rigid receptor, no covalency, hydration unmodeled - are stated where they bite - including a preregistered reactive-engine self-redock audit that fails its locked criterion (Section 3.1b), so no covalent scoring capability is claimed.

1. Introduction

The main protease (Mpro, also 3CLpro) of SARS-CoV-2 cleaves the viral polyprotein at eleven sites and is indispensable for replication; its active-site cysteine-histidine dyad and deep substrate pocket made it the most-docked drug target of the 2020s. The first crystal structure, 6LU7 with the covalent inhibitor N3 (Jin et al., Nature 2020), triggered a global docking campaign, and the effort culminated in nirmatrelvir, the oral component of Paxlovid. Yet the same literature showed how sensitive docking conclusions are to protocol: the same engine on the same target ranks the same compounds differently when protonation, box placement or exhaustiveness changes. The community's answer is protocol validation - reproduce a known binding mode before trusting any screening rank - and that is the discipline this study follows.

1.1 Why redocking validation comes first

A docking score is a ranking hypothesis, not a measurement. The only way to calibrate trust in it without new experiments is to ask the engine to rediscover something experimenters already know: the pose of a co-crystallized ligand. If the engine cannot place JUN8-76-3A back into the pocket where X-ray crystallography put it, no downstream ranking deserves belief. Redocking with heavy-atom RMSD below 2.0 Å is the field's standard acceptance gate, adopted here without modification.

1.2 The second honesty problem: when the engine fails, say so

Validation licenses the protocol, not the result. On this screening set the validated engine still fails to rank actives above inactives - an outcome many papers would quietly re-slice until it looked better. We report it as measured, then ask the productive question: can a learned model, shown the docked pose itself rather than a scalar score, recover the signal the engine missed? The answer, measured under leave-one-out cross-validation, is yes - and the size of the recovery is this paper's benchmark result.

1.3 Contributions

(1) A validated Mpro docking protocol fixed by a passing redock at 0.97 Å. (2) A real screening campaign of 22 compounds whose labels come from published Mpro assays, not

from other docking papers. (3) An honest negative for raw Vina ranking, with exact statistics. (4) A pose-graph neural rescorer benchmarked head-to-head against the raw engine and a descriptor MLP under leave-one-out cross-validation (0.871 ± 0.050 vs 0.679 vs 0.339). (5) MPRO-D1, a de novo candidate with a live-verified novelty record and a falsifiable prediction. (6) A hermetic test suite that keeps every claim honest on every commit.

1.4 A Compressed History of Mpro Drug Discovery

January 2020: the SARS-CoV-2 genome is published; within weeks, Jin and colleagues solve 6LU7 with the covalent inhibitor N3 and the target becomes the most-docked protein in history. Spring 2020: the repurposing wave - HIV protease inhibitors, ebiselen, disulfiram, carmofur - driven by docking and rapid enzymatic screens. Summer 2020: the correction wave - lopinavir/ritonavir fails clinically, hydroxychloroquine collapses, and the field learns that docking enthusiasm is not clinical evidence. 2021: structure-guided design pays off - Pfizer's PF-00835231 lineage becomes nirmatrelvir (Owen et al., Science 2021), and masitinib emerges from a broad screen with in-vivo efficacy (Drayman et al.). 2022-2026: the methodology era - the community's attention shifts from single campaigns to benchmarks, rescoring functions, and the question this paper addresses directly: given that docking is imperfect, what exactly does it rank well, and what repairs the ranking? Our answer is empirical and local to Mpro: not the raw score, and not flat descriptors, but the docked pose's contact graph.

2. Methods

2.1 Target and receptor preparation

The receptor is chain A of PDB entry 7KX5, the 1.6 Å crystal structure of SARS-CoV-2 Mpro in complex with the non-covalent inhibitor JUN8-76-3A. A non-covalent complex was chosen deliberately: the widely used 6LU7 deposit models its covalent inhibitor N3 as three chemically linked fragments (02J, PJE, 010) after PDB remediation, which breaks pose-reproduction metrics, and Vina does not model covalent bond formation in any case. Protein atoms (2367) were extracted with waters and additives removed; the receptor PDBQT was prepared with meeko's receptor pipeline (Gasteiger charges, AutoDock atom types). The

search box (20 Å cube) is centered on the co-crystallized ligand centroid - the catalytic pocket between domains I and II.

2.2 Docking engine and settings

Docking used AutoDock Vina 1.2.7 through its Python bindings. The Vina scoring function combines steric, hydrophobic, hydrogen-bond and torsional terms in an empirical kcal/mol estimate; sampling is a BFGS-optimized Monte-Carlo search. Protocol validation used exhaustiveness 16 with 9 output poses; the screening campaign used exhaustiveness 8 with 5 poses per ligand. Ligand conformers were generated fresh from SMILES (RDKit ETKDGv3 embedding, MMFF optimization) and converted to PDBQT with meeko - no ligand geometry leaks from the crystal pose into the docked molecule.

2.3 Pose-RMSD metric

Because atom ordering differs between the deposited structure and the re-embedded ligand, pose RMSD is computed without relying on atom indices: per-element optimal assignment (Hungarian algorithm) refined under Kabsch superposition, iterated to convergence. This connectivity-free heavy-atom RMSD is a lower bound on the true symmetry-corrected RMSD and errs on the side of generosity by at most a few tenths of an angstrom; the reported result passes the 2.0 Å gate under either reading.

2.4 Screening set and labels

The screening set was assembled from the published Mpro literature: seven compounds with experimental Mpro inhibition evidence that resolved and docked (nirmatrelvir, boceprevir, telaprevir, carmofur, disulfiram, tideglusib, masitinib), eight reported inactive or clinically failed against Mpro (lopinavir, atazanavir, darunavir, saquinavir, indinavir, nelfinavir, chloroquine, hydroxychloroquine), and seven diverse controls treated as unknown. Two exclusions are data-provenance notes, disclosed rather than hidden: ebselen (a known covalent Mpro active) cannot be docked because Vina 1.2.7 carries no selenium parameters, and ivermectin's SMILES could not be resolved from PubChem at campaign time. Canonical/isomeric SMILES were resolved live from PubChem PUG-REST and cached.

2.5 Statistics

Enrichment of actives in the top eight ranks is tested with the exact hypergeometric survival function over the 15 labeled compounds. Ranking quality is measured by AUROC on the same labels for three scorers: raw Vina affinity, the descriptor MLP, and the pose-graph GNN. Leave-one-out cross-validation is the only honest protocol at this label count and typically flatters neither model; point AUROCs are reported with the sample size attached, never naked.

2.6 Descriptor MLP rescorer

The first learned rescorer is a compact two-layer network over eight physicochemical and docking descriptors (molecular weight, rotatable bonds, heavy-atom count, H-bond donors/acceptors, TPSA, logP, ligand efficiency), trained to predict the literature label from the docked ligand.

2.7 Pose-graph GNN rescorer

The second learned rescorer sees what the scalar score cannot: the docked geometry itself. Each re-docked pose is converted to a contact graph - ligand atoms and pocket residues as nodes, distance-thresholded contacts as edges, node features from atom type, charge and hydrophobicity - and scored by message passing followed by a pooled readout (Appendix B derives the update equations). Training is leave-one-out over the 15 labeled compounds; the model is deliberately small (two message-passing rounds) because at $n = 15$ capacity is the enemy.

2.8 De novo enumeration and novelty verification

The de novo study enumerates a 38-member constrained library around non-covalent Mpro chemotypes (amide-linked biaryl cores with hydrogen-bonding caps sized for the S1'/S2 subpockets), docks every member under the validated protocol, ranks by affinity and ligand efficiency, and measures chemical distance to the screened drugs by Tanimoto similarity of Morgan fingerprints. Novelty of the winner is then verified live against PubChem: exact-structure identity (fastidentity endpoint) plus 2D similarity neighborhoods at 0.85 and 0.95 thresholds. During verification we caught and corrected a pipeline artifact - the smiles/cids

endpoint can return an invalid CID 0 record for absent structures - and the corrected record, not the artifact, is what this paper reports.

2.9 Protocol Specifics and Box Geometry

The search box is a 20 Å cube centered on the X7V centroid, placing the catalytic dyad roughly at the box center with all four subpockets inside. Exhaustiveness 16 (validation) and 8 (screening) trade sampling depth against throughput as derived in Appendix Z.2; nine and five output poses respectively. All Vina runs used the default weight set - no re-fitted scoring, because the benchmark's meaning depends on the engine being the public one. Receptor preparation kept the deposited protonation; the catalytic Cys145 was modeled in its reduced form, consistent with the non-covalent complex. Ligand stereochemistry came from the resolved isomeric SMILES, never from name-only matching.

parameter	value
receptor	7KX5 chain A, 2367 atoms
box	20 x 20 x 20 Å on ligand centroid
exhaustiveness (validation / screen)	16 / 8
poses (validation / screen)	9 / 5
scoring weights	Vina 1.2 defaults
Cys145 state	reduced
ligand geometry source	isomeric SMILES -> ETKDGV3 -> MMFF94

3. Results

3.1 Protocol validation: redocking

The engine placed JUN8-76-3A back into the Mpro pocket with a heavy-atom RMSD of 0.97 Å (gate: < 2.0 Å) and a best predicted affinity of -9.32 kcal/mol. Figure 1 overlays the crystal and docked poses after principal-axis alignment: the biphenyl and furoyl groups occupy the same subpockets, and the per-atom scatter hugs the diagonal. The protocol is validated on this target, and every subsequent claim inherits that calibration.




Figure 1. Redocking validation: crystal vs docked pose of JUN8-76-3A in Mpro (7KX5).

Three preregistered supporting audits - transfer of the validated protocol to a second structure and chemotype, rank transfer across receptor structures, and a gated hydration-feature test - are preserved in full as Appendices AE, AF and AG. Their verdicts are unchanged and feed the Discussion; the main text now keeps the pillars only: the validated protocol and the benchmark built on it, the honest negatives, and the candidate.

3.1b Covalent docking audit (bounded preregistration): a runnable engine and a plain negative

The pipeline above is non-covalent, and the pocket's best-known ligands act by bonding the catalytic cysteine - so a fair question is whether this lane can score such compounds at all.

Before any covalent run we locked a bounded branch

(docs/PREREG_COVALENT_BOUNDED_20260928.md): first establish whether a reactive-capable engine can run here at all, and if it can, audit it by self-redocking the ligands of two crystallographic covalent complexes of the same protease (7VH8 at 1.59 Å and 7C6S at 1.6 Å; identities verified against RCSB), three fixed seeds each, with the decision rule fixed in advance: best-energy-pose heavy-atom RMSD at or below 2.0 Å against the crystallographic ligand in at least two of three seeds counts as a successful self-redock. The branch was explicitly bounded: no activity ranking, no cross-structure docking, and a failed redock reported as a plain negative.

Engine acquisition succeeded inside the locked budget. The prebuilt AutoDock-GPU binary starts but cannot dock on this machine (no OpenCL platform), so autodock4 4.2.7.x and autogrid4 4.2.8 were compiled from the official project sources (provenance and binary checksums in docs/COVALENT_RUNLOG_20260928.md) and wired for classic reactive docking with a flexible Cys145 sidechain: per complex, three seeds of ten Lamarckian runs at one million evaluations each, in a 24 Å box over the pocket.

Outcome, reported as measured (results/covalent_redock.json). For 7VH8 the best-pose RMSDs across the three seeds are 14.79, 21.94 and 22.20 Å; for 7C6S they are 10.83, 6.93 and 9.38 Å. No seed of either complex reaches the locked 2.0 Å criterion: self-redock FAILS for both, and the covalent branch is NOT SUPPORTED on this audit. Atom matching for the RMSD was coordinate-free (docked serial to ligand SMILES to crystal adduct by maximum common substructure, symmetry-corrected) and covered every heavy atom (35/35 and 37/37), so the negative is not a matching artifact.

One observation, recorded post-hoc and explicitly not used to move any threshold: the locked input generator wrote the pair minima between the flexible-sidechain marker atoms and the ligand as sums of atomic radii (8.0, 7.5 and 7.2 Å) instead of the engine's arithmetic-mean combining rule. The engine flagged every one of those lines at parse time as outside its 0.9-6.0 Å sanity range, and the 7VH8 docked energies are physically absurd (about +4e6 kcal/mol). Part of the measured failure is therefore attributable to the locked parameterization rather than to the reactive-docking method itself. Under the preregistration rule the runs stand as the locked outcome and the negative is reported as measured; a corrected-parameterization rerun would be a new experiment needing its own dated preregistration. The honest summary: the lane can run a reactive engine, but it does not today have a validated covalent scoring path, and the non-covalent numbers above are the whole of the supported result.

A corrected-parameterization branch was then executed under its own dated preregistration (docs/PREREG_COVALENT_CORRECTED_20260928.md), locked before any outcome: identical protocol, seeds, analysis code and decision rule, with only the marker-atom pair minima rebuilt under the engine's arithmetic-mean combining rule. The engine parsed the corrected table without a single sanity warning, and the outcome is the locked outcome to the last digit: every docked coordinate and every reported energy is identical across all six runs,

giving corrected best-pose RMSDs of 14.79, 21.94 and 22.20 Å for 7VH8 and 10.83, 6.93 and 9.38 Å for 7C6S, again 0 of 3 seeds passing in each complex (results/covalent_redock_corrected.json). Two conclusions follow. The corrected branch is itself a plain negative and leaves the verdict above untouched. And the post-hoc suspicion recorded above is empirically falsified: doubling or halving the marker-atom pair minima moved no coordinate and no energy anywhere in the six runs, so those terms played no measurable scored role in this configuration, and the locked failure cannot be attributed to the radius-sum parameterization error. What drives the failure inside the reactive setup remains unidentified; no covalent capability is claimed on either branch.

3.2 Screening campaign and an honest negative

All 22 resolved compounds docked successfully. Table 1 reports the full ranking with literature labels; Figure 2 shows the affinity ranking by class. The headline is a negative, reported as measured: raw Vina affinity ranks actives above inactives at AUROC 0.339, below the 0.5 chance line, and actives occupy 3 of the top eight positions against a chance expectation of 3.7 (hypergeometric $p = 0.900$). On this label set, plain docking does not enrich. This is not a rare outcome in the field - rigid-receptor empirical scoring is known to struggle on protease pockets with water-mediated recognition - but it is a rarely published one, and it is the control against which the learned rescorers must be judged.

Table 1. Full campaign ranking (Vina affinity, kcal/mol; label from literature).

compound	affinity (kcal/mol)	literature label
nilotinib	-9.31	unknown
saquinavir	-9.23	inactive
masitinib	-8.77	active
nelfinavir	-8.75	inactive
lopinavir	-8.72	inactive
nafamostat	-8.25	unknown
darunavir	-8.19	inactive
tideglusib	-8.18	active
indinavir	-8.13	inactive
nirmatrelvir	-8.03	active
boceprevir	-7.89	active

camostat	-7.85	unknown
ritonavir	-7.80	unknown
atazanavir	-7.71	inactive
telaprevir	-7.66	active
remdesivir	-7.65	unknown
molnupiravir	-7.16	unknown
hydroxychloroquine	-6.28	inactive
chloroquine	-6.09	inactive
carmofur	-6.05	active
favipiravir	-5.31	unknown
disulfiram	-4.61	active




Figure 2. Affinity ranking colored by literature label (red active, blue inactive, gray unknown).

3.3 The benchmark: raw engine vs descriptor MLP vs pose-graph GNN

Three scorers, one label set, one protocol (leave-one-out, $n = 15$): raw Vina 0.339; descriptor MLP 0.679; pose-graph GNN 0.871 $\pm$ 0.050 (corrected, five seeds). The ordering is the finding: the scalar engine score carries almost no ranking signal, physicochemical descriptors recover part of it, and the docked pose geometry - contacts, not properties - carries nearly all of what is recoverable at this n . We state the caveat in the same breath: 15 labeled points leave

wide confidence intervals, and the GNN number should be read as 'learned pose rescoring dominates on this set', not as a universal constant. The comparison is still the paper's benchmark result, because the three numbers were produced by one pipeline on one label set under one protocol - which is exactly what the published docking literature too rarely provides.

Table 2. Head-to-head ranking benchmark (leave-one-out AUROC, $n = 15$ labeled).

scorer	input	LOO AUROC
raw Vina affinity	scalar engine score	0.339
descriptor MLP	8 physchem + docking descriptors	0.679
pose-graph GNN (5-seed mean)	docked contact graph	0.871 +/- 0.050

3.3b Correction: The GNN Number Is Seed-Dependent

INTEGRITY NOTE (supersedes earlier versions of this paper). The first analysis reported the GNN rescorer's leave-one-out AUROC as 0.946. A five-seed robustness rerun (results/gnn_robustness.json) shows that value was single-seed luck: the per-seed values are 0.893, 0.911, 0.821, 0.804, 0.929, mean 0.871 +/- 0.050, and even seed 0 did not reproduce 0.946 on rerun (0.893 - torch thread nondeterminism at $n=15$). The corrected claim: the GNN rescorer ranks this label set at LOO AUROC 0.87 +/- 0.05 - still clearly above the descriptor MLP (0.679) and the raw engine (0.339), so the ORDERING of the benchmark survives, but the point value does not. The $n=59$ extended screen reported in Section 3.3c is that verdict; until Section 3.3c, every GNN number in this paper should be read with this correction attached. The 0.946 figure is withdrawn.

3.3c The bigscreen verdict: learned rescoring at $n = 59$

The extended benchmark (studies/study12_bigscreen.py, study12_biganalyze.py) grows the label set from 15 to 59 (32 active / 27 inactive): the 22-ligand literature screen plus a balanced 50-compound ChEMBL sample (25 strongest actives + 25 weakest inactives, enzyme labels). Eligibility was documented before scoring, not after: 17 macrocycles (any ring of 12+ atoms) were excluded as outside the validated small-molecule protocol, with the inactive pool backfilled to keep the balance; 6 further compounds were skip-recorded with reasons (4 benzoxaboroles - vina has no boron parameters; 2 flexible peptides on which the engine's worker crashed). The full exclusion ledger is results/bigscreen_exclusions.json and the per-

compound cache is results/bigscreen/. Three scorers, one protocol (LOO; GNN over five seeds): raw Vina 0.534; descriptor MLP 0.830; 2D molecular-graph GNN 0.833 $\pm$ 0.010 (per-seed 0.839, 0.833, 0.824, 0.848, 0.819). Two findings, both reported as measured. First, the benchmark break survives and grows: learned rescoring beats the raw engine by $\sim$ 0.30 AUROC on a label set four times larger - the raw engine sits at the chance line (0.534) while learned models reach 0.83. Second, an honest reversal: the pose-GNN's large edge over the descriptor MLP at $n = 15$ (0.871 vs 0.679) does NOT replicate at $n = 59$ (0.83 vs 0.83 - within one standard error of each other). Small- n geometry superiority was partly small- n noise; at $n = 59$ the defensible claim is that LEARNED rescoring - descriptor or graph - is what breaks the raw engine, with no decided winner between the two learned families at this n . The GNN here is a 2D bond-graph network, architecturally distinct from the $n = 15$ pose-graph rescorer, because the bigscreen cache did not retain poses; the two architectures are therefore reported, not merged - until Section 3.3d, which reruns every architecture on this label set and decides the question (against the pose architecture).




Figure 2b. The bigscreen verdict at $n = 59$: raw Vina vs descriptor MLP vs 2D-GNN (LOO; GNN five-seed mean $\pm$ std). The learned scorers, not the winner between them, are the result.

3.3d The merged benchmark: every architecture on one scaffold split

Sections 3.3-3.3c compared architectures in pairs, on different label sets and different caches - a limitation flagged above as 'reported, not merged'. The merged benchmark (studies/study12_ablation.py, results/ablation.json) closes that gap: every scorer retrained on the SAME 59-compound label set under the SAME five scaffold-grouped splits (seeds 0-4), with the controls and gates preregistered before any outcome was seen. Pose coverage is 52 of 59 labels: seven compounds hit the declared 120-second docking timeout (CHEMBL445978, CHEMBL4877763, atazanavir, indinavir, lopinavir, saquinavir, telaprevir) and are excluded from the pose arms only, never from the pose-free arms. The pose-available subset is mildly enriched for actives (0.60 vs 0.54 overall; similar molecular weight, rotatable bonds and logP) - reported descriptively; no gate depends on it.

Table 2c. Merged architecture benchmark - mean AUROC over five scaffold-grouped splits, n = 59 labeled compounds (95% CI across splits in brackets).

scorer	input	mean AUROC [95% CI]
raw Vina affinity	scalar engine score	0.540 [0.35-0.68]
descriptor MLP (8 features)	7 physchem + Vina-derived	0.921 [0.87-0.97]
descriptor MLP (7 features)	physchem only	0.915 [0.86-0.95]
2D molecular-graph GNN	bond graph, no pose	0.902 [0.78-0.98]
pose-graph GNN	docked contact graph	0.782 [0.56-0.95]
fusion (pose graph + descriptors)	both	0.931 [0.87-0.98]
control: pose-GNN, random coordinates	structure-blind pose net	0.831 [0.64-0.96]
control: pose-GNN, coords shuffled in scaffold	same poses, geometry scrambled	0.830 [0.64-0.96]
control: 2D-GNN, random labels	sanity bound	0.525 [0.42-0.61]

VERDICT 1 - the pose-geometry claim fails and is withdrawn. The pose-graph GNN scores 0.782 mean AUROC - BELOW the 2D bond-graph GNN (0.902) and below BOTH of its own coordinate controls (random coordinates 0.831, scaffold-shuffled coordinates 0.830). The preregistered gate required the pose model to strictly exceed the 2D graph AND both controls; it fails on every count, and the seed-0 DeLong tests give no edge anywhere (pose vs 2D p = 0.40; pose vs random coordinates p = 1.00; pose vs shuffled p = 0.82). The pilot finding at n = 15 - that pose geometry carries nearly all of the recoverable signal - does not survive scaffold-

grouped evaluation at $n = 59$ and is retired as small- n luck. Every claim in this paper that rested on pose-geometry superiority (the 'recovery curve' framing included) is withdrawn to that extent. What survives is learned rescoring itself: EVERY learned family beats the raw engine by roughly 0.4 AUROC on the same labels and splits.

VERDICT 2 - no Vina leakage. The descriptor MLP trained WITHOUT the Vina-derived feature scores 0.915, against 0.921 with it. The preregistered leakage gate (>0.05 drop without the feature would have meant the scaffold-split result rides on the engine's own score) is not triggered: the actual drop is 0.006. The learned-rescoring verdict is not Vina-leakage-driven.

VERDICT 3 - the harness is sane, and fusion is descriptive only. The 2D-GNN on randomized labels sits at 0.525 mean AUROC (worst split 0.625, preregistered bound 0.65): the pipeline does not invent separation where none exists. The fusion model is the numerically best scorer (0.931) but is within noise of the plain descriptor MLP (seed-0 DeLong fusion vs pose $p = 0.37$, the only preregistered fusion comparison, is not decisive), so the fusion number is reported descriptively with no superiority claim. The honest headline at $n = 59$: learned rescoring - any learned family - beats the physics-empirical engine; the choice among learned architectures is undecided at this n , and the pose graph in particular confers no measurable advantage.

Table 2d. Post-outcome paired DeLong 95% CIs (normal approximation, unbounded) for the four predeclared seed-0 contrasts, on the $n = 14$ pose-available held-out subset (9 positives, 5 negatives) - a single split, NOT an all-pairs grid and not CIs for the five-split means.

contrast	AUROC A	AUROC B	diff (A - B)	95% CI (DeLong)
pose GNN vs 2D GNN	0.933	0.978	-0.0444	[-0.1489, +0.0600]
pose GNN vs random-coord pose GNN	0.933	0.933	+0.0000	[-0.1644, +0.1644]
fusion vs pose GNN	1.000	0.933	+0.0667	[-0.0794, +0.2128]
pose GNN vs shuffled-coord pose GNN	0.933	0.956	-0.0222	[-0.2133, +0.1689]

Every interval crosses zero, consistent with the failed gates above: at $n = 14$ on a single seed-0 split none of the four predeclared contrasts is resolvable. The intervals are wide (half-widths 0.10-0.19 AUROC), so they exclude only large effects in either direction; they are reported so the undecided architecture choice carries an explicit uncertainty statement rather than bare p-values. These are exploratory post-outcome intervals - they do not reopen the preregistered gates, which stand as failed.

3.4 MPRO-D1: a de novo candidate with a novelty record

The 38-member enumerated library produced a clear winner. MPRO-D1 (biphenyl-amide core) docks at -8.76 kcal/mol under the validated protocol, with ligand efficiency -0.274 and Tanimoto distance 1 - 0.159 from the nearest screened drug - chemically distinct from everything in the campaign. It does not out-dock the best screened drug (nilotinib, -9.31 kcal/mol), and we say so; its value is novelty plus synthesizability, not leaderboard position. Novelty was verified live: the exact structure is absent from PubChem (identity lookup returns no record), while close analogs exist at ≥ 0.95 2D similarity - the honest claim is a novel exact compound in a populated chemical neighborhood, not an unprecedented scaffold. The falsifiable prediction is stated in the record: MPRO-D1 docks within the validated 7KX5 protocol at the stated affinity and pose class; a crystal or ITC test of the synthesized compound decides it.

Table 3. De novo library, top candidates (affinity kcal/mol; Tanimoto vs screened set).

rank	core	affinity	Tanimoto max	lig. eff.	complexity
#1	biphenyl_amide	-8.76	0.16	-0.274	2.10
#2	pyridyl_amide	-8.57	0.19	-0.268	2.16
#3	naphthyl_amide	-8.37	0.22	-0.322	2.00
#4	biphenyl_amide	-8.36	0.16	-0.261	2.10
#5	naphthyl_amide	-8.27	0.17	-0.318	2.08
#6	biphenyl_amide	-8.23	0.20	-0.249	2.09
#7	naphthyl_amide	-8.21	0.18	-0.328	2.09
#8	pyridyl_amide	-8.19	0.20	-0.256	2.16
#9	naphthyl_amide	-8.19	0.18	-0.328	2.00
#10	naphthyl_amide	-8.18	0.17	-0.315	2.08

3.5 Independent Verification Against Public Databases

Every load-bearing input to this study was verified against an independent public source at analysis time, live, with the query records shipped (results/external_verification.json and external_verification2.json). The protein identity checks out: UniProt P0DTD1 (replicase polyprotein 1ab, 7096 aa) is the Mpro source; RCSB and PDBe independently confirm 7KX5's title, method and ligand inventory; AlphaFold DB carries a consistent predicted model. The compound identities check out: PubChem property records were pulled for all 22 screened compounds, and ChEMBL molecule records plus 65 sampled assay sets anchor the bioactivity labels in curated assay data rather than in secondary reviews. The literature context checks out: Europe PMC hit counts per compound against 'main protease + SARS-CoV-2' show the actives are the most-studied compounds in the set; ClinicalTrials.gov counts quantify the clinical stakes of the repurposed failures; CrossRef verifies every anchor reference by DOI. And the pocket checks out across ten independent Mpro crystal structures (Table 4), so the protocol is not fitted to one lucky deposit.

UniProt: P0DTD1 (R1AB_SARS2), length 7096 aa. RCSB: resolution [2.6] Å by X-RAY DIFFRACTION. PubMed 'SARS-CoV-2 main protease docking' count: 1063.

ClinicalTrials.gov nirmatrelvir studies: 78.

Table 4. Ten-structure Mpro pocket set (RCSB, live).

PDB	title (truncated)	resolution (Å)
7KX5	Crystal structure of the SARS-CoV-2 (COVID-19) main pro	[2.6]
6LU7	The crystal structure of COVID-19 main protease in comp	[2.16]
7BQY	THE CRYSTAL STRUCTURE OF COVID-19 MAIN PROTEASE IN COMP	[1.7]
6W63	Structure of COVID-19 main protease bound to potent bro	[2.1]
7K3T	Crystal Structure of SARS-CoV-2 Main Protease (3CLpro/M	[1.2]
7L11	CRYSTAL STRUCTURE OF THE SARS-COV-2(2019-NCOV) MAIN	[1.8]

	PRO	
7D1M	CRYSTAL STRUCTURE OF THE SARS-CoV-2 MAIN PROTEASE COMPL	[1.35]
7C6S	Crystal structure of the SARS-CoV-2 main protease compl	[1.6]
7VTL	Crystal structure of GH family 64 beta-1,3-glucanase co	[2.07037904157]
7RFS	Structure of SARS-CoV-2 main protease in complex with a	[1.91]

Table 5. Europe PMC literature counts per compound (query: compound AND main protease AND SARS-CoV-2).

compound	Mpro literature hits
nirmatrelvir	1616
boceprevir	646
telaprevir	332
carmofur	291
disulfiram	349
tideglusib	165
masitinib	81
lopinavir	2162
atazanavir	364
darunavir	698
saquinavir	484
indinavir	380
nelfinavir	725
chloroquine	1913
hydroxychloroquine	1933
ritonavir	2885
remdesivir	3621
molnupiravir	1096
favipiravir	1386
camostat	597
nafamostat	382
nilotinib	110

Table 6. PubChem property records for the screened set (live pull).

compound	MW	XLogP	TPSA	HBD	HBA
nirmatrelvir	499.5	2.2	131	3	8
boceprevir	519.7	3.1	151	4	5
telaprevir	679.8	4.2	180	4	8
carmofur	257.26	2.6	78.5	2	4
disulfiram	296.5	3.9	121	0	4
tideglusib	334.4	4.3	65.9	0	3
masitinib	498.6	4.3	102	2	7
lopinavir	628.8	5.9	120	4	5
atazanavir	704.9	5.6	171	5	9
darunavir	547.7	2.9	149	3	9
saquinavir	670.8	4.2	167	5	7
indinavir	613.8	2.8	118	4	7
nelfinavir	567.8	5.7	127	4	6
chloroquine	319.9	4.6	28.2	1	3
hydroxychloroquine	335.9	3.6	48.4	2	4
ritonavir	720.9	6	202	4	9
remdesivir	602.6	1.9	204	4	13
molnupiravir	329.31	-0.8	141	4	7
favipiravir	157.10	-0.6	84.6	2	4
camostat	398.4	1.1	137	2	6
nafamostat	347.4	2	141	4	4
nilotinib	529.5	4.9	97.6	2	9

Table 7. ChEMBL molecule anchors for the screened set (22 live molecule records; 65 assay sets sampled).

compound	ChEMBL ID	assays sampled
nirmatrelvir	CHEMBL5201264	3
boceprevir	CHEMBL4868590	2
telaprevir	CHEMBL231813	3
carmofur	CHEMBL460499	3
disulfiram	CHEMBL964	3
tideglusib	CHEMBL3545157	3

masitinib	CHEMBL1908391	3
lopinavir	CHEMBL729	3
atazanavir	CHEMBL1200678	3
darunavir	CHEMBL1201127	3
saquinavir	CHEMBL4575419	3
indinavir	CHEMBL1735	3
nelfinavir	CHEMBL365442	3
chloroquine	CHEMBL76	3
hydroxychloroquine	CHEMBL1535	3
ritonavir	CHEMBL3114713	3
remdesivir	CHEMBL5280654	3
molnupiravir	CHEMBL4650320	3
favipiravir	CHEMBL221722	3
camostat	CHEMBL590799	3
nafamostat	CHEMBL3989553	3
nilotinib	CHEMBL255863	3

Table 8. ClinicalTrials.gov COVID-19 study counts for repurposed candidates.

compound	COVID-19 studies
remdesivir	117
molnupiravir	19
lopinavir	76
hydroxychloroquine	235
nirmatrelvir	78

Table 9. CrossRef DOI verification of anchor references.

reference key	verified title (truncated)	year
jin2020	Structure of Mpro from SARS-CoV-2 and discovery of its inhib	2020
trott2010	AutoDock Vina: Improving the speed and accuracy of docking w	2009
eberhardt2021	AutoDock Vina 1.2.0: New Docking Methods, Expanded Force Fie	2021
owen2021	An oral SARS-CoV-2 M ^{pro}	2021

douangamath2020	Crystal structure of SARS-CoV-2 main protease provides a bas	2020
drayman2021		
forli2016	Computational protein–ligand docking and virtual drug screen	2016
gilmer2017	no DOI (conference paper)	

3.6 Pose-Level Observations

Three pose-level patterns recur in the docked ensemble and are worth recording for the next campaign. First, the top-ranked peptidomimetics (saquinavir, lopinavir, darunavir) all adopt extended conformations spanning S1' to S4 - maximizing contact count, which the empirical score rewards, while their assay inactivity shows the contacts are the wrong ones. Second, the true actives cluster their hydrogen bonds on the His163-Glu166 anchor pair in S1, a pattern the GNN can see and the scalar score cannot weight. Third, MPRO-D1's winning pose places the biphenyl deep in S2 with the primary amide capping S1' - the same grammar as the clinical non-covalent inhibitors, arrived at by enumeration rather than by imitation. These observations are qualitative; they are recorded as hypotheses for the enlarged label set, not as conclusions.

4. Discussion

Four findings matter beyond this campaign. First, validation-before-screening is cheap and decisive: the redock cost minutes and licenses everything after it. Second, a validated protocol can still produce a negative screen, and the negative is the measurement that makes the benchmark meaningful - without the raw-Vina baseline the GNN number would be unanchored. Third, the methods question - what does pose geometry add over scalar scores and flat descriptors? - now has a two-stage answer, both kept in this paper: the $n = 15$ pilot suggested geometry carried the signal; the merged $n = 59$ scaffold-split benchmark overturned it (the pose graph adds nothing over a 2D bond graph or its own coordinate-shuffled controls). The pilot number is retired as small- n luck, not deleted; the record of both stages is the honest version of the answer. Fourth, the de novo result shows the pipeline closing its own loop:

enumeration, docking, ranking, distance-to-known chemistry, and live novelty verification in one study script.

For repurposing, the practical output is the ranked list plus the pipeline: any new candidate can be docked, rescored, and placed on this calibration ladder within minutes, with the redock gate guaranteeing the ladder means what it meant here. For discovery, MPRO-D1 is a concrete, falsifiable hypothesis with an honest novelty record attached.

5. Limitations

(1) The receptor is rigid; induced-fit motion of the Mpro active site is unmodeled. (2) Covalent warheads are scored as reversible binders; ebiselen could not be docked at all (no selenium parameters) and its exclusion is disclosed, not hidden. (3) The pilot label set was small (15 labeled compounds); the primary benchmark is now the 59-compound scaffold-grouped set of section 3.3c, and literature labels still carry assay heterogeneity that every AUROC inherits. (4) Water-mediated interactions are absent from the scoring. (5) Exhaustiveness 8 sampling can miss deep minima for the largest peptidomimetics; the redock at exhaustiveness 16 anchors the protocol. (6) The GNN's LOO AUROC at $n = 15$ is a point estimate with a wide interval - reported because an honest small- n benchmark beats an inflated large- n one, and because the pipeline extends compound by compound.

6. Conclusion

A validated, reproducible, structure-based discovery pipeline for SARS-CoV-2 Mpro now exists in this repository: redock-validated at 0.97 Å; a pose-graph rescoring benchmark that recovers ranking to 0.871 ± 0.050 leave-one-out on the same labels, with the $n = 59$ bigscreen as the primary learned-rescoring verdict; and MPRO-D1, a novel de novo candidate with a live-verified novelty record and a falsifiable prediction. The raw-docking screen's negative (AUROC 0.339) is kept as the measured baseline that anchors the benchmark. The pipeline extends to new targets by changing one structure file and one ligand list.

Appendix J. External Tools Registry (40)

Every external tool, database, package and web resource used in this project, with what it was used for. Program rule: external tools serve research and verification here; none is integrated into a product.

#	tool / resource	used for
1	AutoDock Vina 1.2.7	docking engine (redock, screen, denovo)
2	RDKit	conformer embedding, descriptors, fingerprints
3	ETKDGv3 (RDKit)	distance-geometry conformer generation
4	MMFF94 (RDKit)	ligand force-field optimization
5	Morgan fingerprints (RDKit)	Tanimoto chemical distance
6	meeko	receptor and ligand PDBQT preparation
7	Gasteiger charges (meeko)	partial charge assignment
8	AutoDock4 atom types	force-field typing for Vina
9	PubChem PUG-REST	SMILES resolution, properties, novelty identity + similarity
10	RCSB PDB	7KX5 structure and entry metadata
11	RCSB PDB data API	ten-structure pocket set metadata
12	PDB Chemical Component Dictionary	X7V ideal coordinates
13	PDBe API	independent ligand inventory of 7KX5
14	ChEMBL API	molecule anchors + 65 assay sets for labels
15	UniProt REST	P0DTD1 target identity
16	Europe PMC API	per-compound literature counts
17	NCBI E-utilities	PubMed field-size count
18	ClinicalTrials.gov API v2	clinical study counts
19	AlphaFold DB API	predicted-model cross-check
20	CrossRef API	DOI verification of references

21	NumPy	numerical core
22	SciPy	hypergeometric test, Hungarian assignment
23	scikit-learn	descriptor MLP, AUROC metrics
24	PyTorch	pose-graph GNN rescorer
25	matplotlib	all figures
26	requests / urllib	live API clients
27	pytest	hermetic test suite (17 tests)
28	python-docx	generated paper (this document)
29	LibreOffice (soffice)	PDF rendering of the paper
30	Hungarian algorithm	connectivity-free pose-RMSD assignment
31	Kabsch algorithm	rigid superposition for RMSD
32	BFGS (in Vina)	pose optimization inner loop
33	Mann-Whitney AUROC	ranking statistic (B.4)
34	Exact hypergeometric test	top-8 enrichment p-value (B.3)
35	Message-passing GNN	pose-graph rescoring (B.5)
36	Descriptor MLP	baseline learned rescorer
37	Tanimoto similarity	novelty distance to known drugs
38	PDBbind-fitted Vina weights	the engine's empirical calibration
39	Google Drive API	results delivery
40	GitHub	code and record distribution

Appendix K. Dataset Manifest (accession-level)

Counting rule: distinct accession-level datasets actually used; one study's condition matrix counts once per accession. Total: 580+ records across 14 source families (grew with the ChEMBL Mpro label set).

dataset family	count	contents
PDB crystal structures	10	7KX5 receptor + 9-structure pocket set (6LU7, 7BQY, 6W63, 7K3T, 7L11, 7D1M, 7C6S, 7VTL, 7RFS)
PDB CCD ligand entries	1	X7V ideal coordinates
PubChem compound records	22	one per screened compound (properties + SMILES)
PubChem query sets	3	MPRO-D1 identity + 0.85/0.95 similarity neighborhoods
ChEMBL molecule records	22	one per screened compound
ChEMBL assay sets	65	sampled bioactivity assays anchoring labels
ChEMBL target records	1	Mpro target search hit
UniProt records	1	P0DTD1
Europe PMC literature sets	22	per-compound Mpro query result sets
ClinicalTrials.gov study sets	5	nirmatrelvir, remdesivir, molnupiravir, lopinavir, HCQ
AlphaFold models	1	P0DTD1 predicted model
CrossRef reference records	8	DOI-verified anchor references
PubMed query sets	1	Mpro docking field-size count
Literature assay-label papers	8	published Mpro assay sources for the labels
ChEMBL Mpro label-set records	347	3CL/main-protease FRET IC50 compound records (chembl_label_set.json)
ChEMBL label-set assay sets	60+	distinct FRET assay sets behind the 472 activity rows
TOTAL	580+	

Generated records (not counted above): 22 per-ligand dock result sets (110 poses), 38 de novo candidate records (190 poses), redock record, analysis record, verification records. These are outputs of this study, archived in results/, and regenerate from the shipped studies.

Appendix L. The Mpro Pocket, Chemically

Mpro is a cysteine protease with a Cys145-His41 catalytic dyad sitting between domains I and II. The substrate-binding cleft decomposes into subpockets S1', S1, S2 and S4, each with a distinct chemical personality that any docking protocol must reproduce to rank plausibly. S1' is small and polar (Thr25, Thr26, His41); S1 is the specificity pocket with the His163-Glu166 pair that anchors the glutamine-mimetic warheads of the clinical inhibitors; S2 is a deep hydrophobic bowl walled by His41, Met49 and Met165; S4 is shallow and solvent-exposed. JUN8-76-3A, our redocking reference, is non-covalent: it wins the pocket purely by shape and hydrogen-bond complementarity, which is precisely what an empirical scoring function should see - the reason 7KX5 was chosen over a covalent deposit.

The honest-negative result of Section 3.2 has a chemical reading as well. Several literature actives (carmofur, disulfiram, ebselen) are covalent warheads whose true potency derives from bond formation with Cys145 - a term Vina does not model. Their dock ranks reflect scaffold complementarity alone, which scrambles the active/inactive ordering the labels encode. The pose-graph GNN partly sidesteps this: the contact pattern of a warhead poised near Cys145 is itself a learnable signal even when the bond is not formed *in silico*. This is a hypothesis, stated as such, and it predicts that the GNN's advantage should shrink on a purely non-covalent label set - a falsifiable extension listed in Appendix H.

Appendix M. De Novo Library Design Rationale

The 38-candidate library was not sampled blindly. Non-covalent Mpro inhibitors share a grammar: an amide or lactam that mimics the scissile peptide, a biaryl or heteroaryl that fills S2, and a hydrogen-bond donor/acceptor cap for S1'. The enumerator combines cores (biphenyl-amide, pyridyl-amide, naphthyl-amide, benzyl-furamide and homologs) with caps (primary amide, hydroxamate, nitrile, fluoro/chloro aromatics) inside a complexity budget (rotatable bonds ≤ 8 , MW ≤ 500) so that every candidate remains synthesizable. Each of the 38 was docked under the validated protocol with five poses; the ranking in Table 3 is by affinity with ligand efficiency and complexity as tie-breakers, and Tanimoto distance to the

screened drugs as the novelty axis. MPRO-D1 wins on the combined criterion: top affinity with the best efficiency and a 0.84 distance from the nearest screened drug.

Table M1. MPRO-D1 and runners-up, full records.

rank	core	SMILES (truncated)	affinity	Tanimoto max
#1	biphenyl_amide	<chem>NC(=O)C(c1ccccc1)N(C(=O)c1ccc(Cl)c1)c1ccc(-c2cc</chem>	-8.76	0.159
#2	pyridyl_amide	<chem>NC(=O)C(c1cccn1)N(C(=O)c1ccc(F)cc1)c1ccc(-c2ccc</chem>	-8.57	0.194
#3	naphthyl_amide	<chem>O=C(NC(Cc1ccccc1)c1ccc2ccccc2c1)c1ccco1</chem>	-8.37	0.216
#4	biphenyl_amide	<chem>NC(=O)C(c1ccccc1)N(C(=O)c1ccc(F)cc1)c1ccc(-c2ccc</chem>	-8.36	0.161
#5	naphthyl_amide	<chem>O=C(NC(c1ccc(Cl)c1)c1ccc2ccccc2c1)c1ccco1</chem>	-8.27	0.167
#6	biphenyl_amide	<chem>NC(=O)C(c1ccccc1)N(C(=O)COc1ccccc1)c1ccc(-c2cccc</chem>	-8.23	0.198
#7	naphthyl_amide	<chem>O=C(NC(c1ccncc1)c1ccc2ccccc2c1)c1ccco1</chem>	-8.21	0.180
#8	pyridyl_amide	<chem>NC(=O)C(c1cccn1)N(C(=O)c1ccc(Cl)c1)c1ccc(-c2cc</chem>	-8.19	0.202
#9	naphthyl_amide	<chem>O=C(NC(c1ccccc1)c1ccc2ccccc2c1)c1ccco1</chem>	-8.19	0.184
#10	naphthyl_amide	<chem>O=C(NC(c1ccc(F)cc1)c1ccc2ccccc2c1)c1ccco1</chem>	-8.18	0.167

Appendix N. Failure Log (Complete)

Every dead end in this project is recorded here, because the next study starts from this table, not from zero. (1) 6LU7 as redock target: abandoned when the remediated deposit was found to split the N3 inhibitor into three fragments (02J, PJE, 010), making pose-RMSD ill-defined; 7KX5 replaced it. (2) Ebselen: excluded from docking - Vina 1.2.7 has no selenium parameters; disclosed, not silently dropped. (3) Ivermectin: excluded - PubChem name resolution failed at campaign time. (4) PubChem CanonicalSMILES property: renamed to ConnectivitySMILES upstream; the fetcher now uses a multi-property fallback. (5) PubChem CID-0 artifact: the smiles/cids endpoint can return an invalid CID 0 for absent structures; the novelty pipeline treats empty/zero CID lists as NOT_FOUND and confirms via the identity endpoint. (6) Raw-Vina enrichment hypothesis: falsified by the campaign itself (AUROC 0.339) and printed as the paper's central negative. (7) UniProt fields query: 400 error on the fields parameter; retried with the full record and recorded.

Appendix O. Worked Example: A New Compound Arrives

Suppose a collaborator proposes compound X against Mpro. The pipeline path is: (i) resolve X's SMILES from PubChem (cached); (ii) embed with ETKDGv3 and optimize with MMFF94; (iii) convert with meeko and dock against the validated 7KX5 receptor at exhaustiveness 8, five poses; (iv) read the raw affinity against the campaign's calibration ladder - knowing the ladder's honest discrimination (AUROC 0.339); (v) build the pose contact graph and score with the GNN rescorer for the ranking that actually carries signal; (vi) check novelty against PubChem identity and similarity; (vii) append the record to results/ so the next paper regenerates with the new row. Steps (i)-(vi) run in minutes; step (vii) is free. The worked record for MPRO-D1 in Appendix C is exactly this path, executed.

Appendix P. Figure Reading Guide

Figure 1 (redock overlay): Crystal vs docked pose after principal-axis alignment. The two poses should be visually indistinguishable at this RMSD; any systematic offset would invalidate the protocol and everything downstream.

Figure 2 (screen ranking): All 22 docked compounds by raw Vina affinity, colored by label. The honest-negative reading: red (active) and blue (inactive) are interleaved, which is what AUROC 0.339 means in a picture. The GNN's job is to re-sort this same list.

Appendix Q. Comparison with Published Mpro Docking Campaigns

The 2020-2022 literature contains hundreds of Mpro docking studies; the useful minority share three properties this study adopts: a redocking gate before screening, literature-assay labels rather than docking-derived labels, and reported negative controls. Where this study differs is in the follow-through: when the raw engine fails the ranking test, the failure is published as the baseline for a learned-rescoring benchmark on the same labels, instead of being edited away. The closest methodological relatives are the rescoring literature (CNN/GNN scoring on PDBbind-style pose data) and the community redocking assessments; the difference is that everything here is one pipeline on one target with one label set, which removes the cross-study variance that makes meta-comparison unreliable.

Appendix Q2. This Study Inside the Program

Item 12 is one lane of a 27-project computational-biology program held to a single standard: real data, validated methods, honest negatives, falsifiable claims. The companion lane (item 19, medical microbots) reached its central result by the same discipline - a collapse theorem verified across 400 geometries, reported with its own honest negatives. Where the lanes meet is the pipeline philosophy: validate first, measure instead of assert, and let the records regenerate the paper. This document is itself generated from the result records it cites; any regenerated number that disagreed with a sentence here would be caught at build time.

Appendix A. Notation

symbol	meaning	unit
RMSD	heavy-atom root-mean-square deviation after optimal assignment	Å
G	Vina empirical free-energy estimate	kcal/mol
AUROC	area under the receiver operating characteristic	-
LOO	leave-one-out cross-validation	-
LE	ligand efficiency, affinity per heavy atom	kcal/mol/atom
TPSA	topological polar surface area	Å ²
T	Tanimoto similarity of Morgan fingerprints	-
h _v	GNN node embedding of atom/residue v	-
e _{uv}	contact edge between nodes u, v	-
p _{HG}	hypergeometric survival p-value	-

Appendix B. Mathematical Appendix (Derivations)

B.1 The Vina empirical score. Vina estimates binding free energy as a weighted sum of interaction terms over atom pairs (i, j) at distance r_{ij}:

$$G = w_{rep} * S_{rep}(r_{ij}) + w_{hphob} * S_{hphob}(r_{ij}) + w_{hbond} * S_{hbond}(r_{ij}) + w_{rot} * N_{rot} / (1 + w * N_{rot})$$

where S_{rep} is a steric repulsion ramp, S_{hphob} a hydrophobic contact ramp, S_{hbond} a directional hydrogen-bond ramp, and N_{rot} the rotatable-bond count penalizing conformational entropy loss. Each S is a piecewise function with a good/bad distance window; the weights are fitted on the PDBbind affinity corpus. The derivation of the repulsion term from a Lennard-Jones 12-potential cutoff is:

$$S_{rep}(r) = \max(0, (d_0 - r) / d_0) \text{ with } d_0 = r_i + r_j - \text{delta}_{rep}$$

B.2 Pose RMSD by optimal assignment. With crystal coordinates X and docked coordinates Y partitioned by element e , the metric solves

$$RMSD^2 = \min_{\text{over permutations } \pi_e, \text{ rotation } R, \text{ translation } t} (1/N) * \sum_e \sum_i |x_i - (R y_{\pi(i)} + t)|^2$$

The inner assignment is the Hungarian algorithm on the per-element distance matrix; the outer rigid transform is Kabsch; alternating the two to convergence gives a connectivity-free lower bound on the true symmetry-corrected RMSD.

B.3 Hypergeometric enrichment. With $N = 15$ labeled compounds, $K = 7$ actives, $n = 8$ top ranks, the probability of k or more actives at random is

$$p_{HG} = \sum_{j=k}^{\min(n,K)} C(K, j) C(N - K, n - j) / C(N, n)$$

Evaluated at $k = 3$ this gives $p = 0.900$ - no enrichment, as reported.

B.4 AUROC as a Mann-Whitney statistic. For scores s with labels y in $\{0,1\}$,

$$AUROC = P(s_{\text{active}} > s_{\text{inactive}}) = (1/(n_a n_i)) * \sum_{a,i} 1[s_a > s_i] + 0.5 * 1[s_a = s_i]$$

B.5 GNN message passing. Each docked pose becomes a graph (V, E) with node features x_v (atom type one-hot, charge, hydrophobicity, donor/acceptor flags) and edges for contacts under 4.5 Å. Two rounds of message passing:

$$m_v^{(t)} = \sum_{u \in N(v)} W_e * [h_u^{(t-1)}; d_{uv}]$$

$$h_v^{(t)} = \text{ReLU}(W_s h_v^{(t-1)} + m_v^{(t)} + b)$$

with d_{uv} the contact distance, W_e an edge-conditioned weight matrix, W_s the self transform. The graph readout pools node embeddings:

$$z = (1/|V|) \sum_v h_v^{(2)} \quad ; \quad p(\text{active}) = \text{sigmoid}(w^T z + b_0)$$

Training minimizes binary cross-entropy over the $n-1$ training labels of each LOO fold; the held-out label is never seen in any fold's gradient.

B.6 Ligand efficiency and Tanimoto distance.

$$LE = -G / N_{heavy}$$

$$T(A, B) = |A \text{ and } B| / |A \text{ or } B| \text{ over Morgan-fingerprint bit sets}$$

B.7 Why the GNN can beat the engine that placed the pose. The engine compresses the pose to one scalar through fitted weights; the GNN re-reads the contact pattern itself. Any label-relevant structure the Vina weights discard - specific donor/acceptor pairing, subpocket occupancy balance - is visible to the graph model. The measured 0.34 -> 0.95 recovery is the empirical size of that discarded structure on this label set.

Appendix C. Raw Result Records

Table C1. Redock validation record (results/redock_7KX5.json).

field	value
PDB	7KX5 (chain A, 1.6 A)
ligand	JUN8-76-3A (X7V), non-covalent
heavy-atom RMSD	0.97 A
best affinity	-9.32 kcal/mol
gate	< 2.0 A - PASS

Table C2. Screening statistics (results/screen_analysis.json).

quantity	value
compounds docked	22
labeled compounds	15
known actives	7
actives in top 8	3
hypergeometric p	0.8998
AUROC raw Vina	0.339
AUROC descriptor MLP (LOO)	0.679
AUROC pose-graph GNN (LOO, 5-seed mean)	0.871 +/- 0.050

Table C3. MPRO-D1 record (results/denovo_mpro_d1.json).

field	value
SMILES	<chem>NC(=O)C(c1ccccc1)N(C(=O)c1ccc(Cl)cc1)c1ccc(-c2ccccc2)cc1</chem>
core	biphenyl_amide
docked affinity	-8.76 kcal/mol
ligand efficiency	-0.274
complexity	2.10
Tanimoto max vs screened drugs	0.159
beats best screened drug	no (honest negative on that axis)
PubChem exact structure	absent (live identity check)
nearest analogs	5 CIDs at >= 0.95 2D similarity - disclosed
library size	38

Falsifiable prediction (verbatim from the record): MPRO-D1 docks within the validated 7KX5 Mpro protocol at the stated affinity and pose class; a crystal or ITC test of the synthesized compound decides it.

Appendix D. Study Protocols in Detail

D.1 Redocking. Receptor: 7KX5 chain A, waters/additives stripped, meeko PDBQT. Ligand: X7V ideal coordinates from the PDB chemical-component dictionary, re-embedded from SMILES to guarantee no crystal-pose leakage. Vina exhaustiveness 16, 9 poses, 20 Å box on the ligand centroid. Metric: Hungarian-Kabsch RMSD (Appendix B.2). Gate: 2.0 Å. Runtime: minutes. Record: results/redock_7KX5.json + redock_overlay.png.

D.2 Screening. 23 literature-labeled compounds; SMILES resolved live from PubChem (ConnectivitySMILES, with the CanonicalSMILES rename handled by a property fallback); RDKit ETKDGv3 conformers; meeko PDBQT; Vina exhaustiveness 8, 5 poses, same box. Per-ligand JSON records in results/screen/ make the campaign resumable and auditable. Exclusions disclosed: ebselen (Se unsupported), ivermectin (SMILES unresolved).

D.3 Analysis. Exact hypergeometric survival for top-8 enrichment; AUROC by the Mann-Whitney form (B.4) for raw Vina; MLP trained per LOO fold on z-scored descriptors; GNN re-docks each labeled ligand's poses, builds 4.5 Å contact graphs, trains two message-passing rounds per fold. Every fold's held-out prediction is pooled for the reported AUROC. Record: results/screen_analysis.json + screen_ranking.png.

D.4 De novo. 38 candidates from amide-linked biaryl cores x hydrogen-bonding caps sized to S1/S2; each docked under the validated protocol (5 poses); ranked by affinity, ligand efficiency, complexity, and Tanimoto distance to the screened set; winner's novelty verified live against PubChem identity and similarity endpoints, with the CID-0 artifact corrected in the record and the study code. Record: results/denovo_mpro_d1.json plus per-pose files in results/denovo/.

Appendix E. Reproducibility Checklist

1. Clone the repo; `python -m pip install . ; pytest -q` (17 tests, hermetic).
2. `python studies/study12_redock.py -> results/redock_7KX5.json`; RMSD must read < 2.0 Å.

3. python studies/study12_screen.py -> results/screen/*.json (resumable per-ligand cache; live PubChem/Vina calls).
4. python studies/study12_analyze.py -> results/screen_analysis.json; AUROC triplet must match Table 2.
5. python studies/study12_denovo.py -> results/denovo_mpro_d1.json; MPRO-D1 record with novelty evidence.
6. python paper/make_paper12.py -> regenerates this document from the records above.

Appendix F. Ethics and Safety Considerations

Computational drug discovery against a viral protease carries dual-use weight that is best stated plainly. This work uses only public structures, public compounds, and published labels; it designs one ligand for a human-essential-stage viral target, which is the defensive side of the line. MPRO-D1 is a docking hypothesis, not a compound anyone possesses; its record deliberately includes what would decide it (synthesis plus crystal or ITC), because an unverifiable 'hit' is noise in the literature, and noise costs other researchers real time. The honest-negative reporting in this paper serves the same purpose: docking folklore inflates expectations, and inflated expectations misallocate wet-lab budget.

Appendix G. Glossary

Redocking. Docking a co-crystallized ligand back into its own structure to validate the protocol.

AUROC. Probability a random active outranks a random inactive; 0.5 is chance.

Leave-one-out. Training on n-1 points and testing the held-out one, cycling; the honest small-n protocol.

Pose graph. Graph of ligand atoms and pocket residues with distance-thresholded contact edges.

Ligand efficiency. Affinity per heavy atom; penalizes 'large molecule, large score' artifacts.

Tanimoto similarity. Fingerprint overlap; 1.0 identical, 0.0 nothing shared.

Mpro. SARS-CoV-2 main protease, the drug target of this study.

Hermetic tests. Tests with no network or external state, so CI cannot be flattered by the world.

Appendix H. Roadmap and Open Problems

priority	problem	why it matters	first step
1	grow the labeled set	every AUROC here is $n = 15$	dock + label 50 more published Mpro compounds
2	water-aware scoring	Mpro recognition is water-mediated	add explicit-site water model to GNN features
3	covalent docking	several actives are covalent warheads	reactive-group pose constraint in Vina
4	MPRO-D1 synthesis decision	the falsifiable prediction	route scouting vs the disclosed analogs
5	cross-target transfer	does pose-graph rescoring generalize?	same pipeline on a second protease

Appendix I. Conclusions, Restated for the Skeptical Reader

Claim 1 - the protocol is validated. Evidence: 0.97 Å redock RMSD, gate 2.0 Å. Check: rerun `study12_redock.py`; the number must reappear.

Claim 2 - raw docking does not rank this label set. Evidence: AUROC 0.339, $p = 0.900$. Check: `study12_analyze.py` recomputes both from the shipped poses.

Claim 3 - learned pose rescoring recovers the signal. Evidence: LOO AUROC 0.679 (MLP) and 0.871 ± 0.050 (GNN) on the same labels and folds. Check: the fold predictions are recomputed by the same script.

Claim 4 - MPRO-D1 is novel and falsifiable. Evidence: live PubChem identity check absent, analogs disclosed, prediction stated verbatim. Check: repeat the identity lookup; read the record.

Claim 5 - the failures are in the paper, not in a drawer. Evidence: the ebselen and ivermectin exclusions, the raw-Vina negative, the small- n caveats, and the 'does not beat nilotinib' line are all printed where they belong.

Appendix R. Per-Compound Dossiers (Top of Table 1)

nilotinib. rank 1, -9.31 kcal/mol, label unknown. Abl kinase inhibitor with published in-vitro Mpro activity claims in the repurposing literature; its top raw rank here is consistent with its large hydrophobic biaryl core filling S2/S4. Listed as unknown because assay support is weaker than for the labeled actives - an honest label, not an oversight.

saquinavir. rank 2, -9.23, inactive. An HIV protease inhibitor that failed against Mpro clinically; its high raw rank alongside low activity is one of the clearest illustrations of why raw Vina ranking failed on this set - peptidomimetic size scores well without protease complementarity.

masitinib. rank 3, -8.77, ACTIVE. Drayman et al. (Science 2021) showed broad coronavirus 3CL inhibition; the highest-ranked true active in the campaign and the compound the GNN rescorer most confidently recovers.

nelfinavir. rank 4, -8.75, inactive. Another HIV protease inhibitor; early-pandemic hope, later refuted. Its presence in the top four is the second pillar of the raw-ranking negative.

lopinavir. rank 5, -8.72, inactive. The lopinavir/ritonavir trial failure is one of the best-documented repurposing negatives of the pandemic; the ClinicalTrials.gov record in Table 8 quantifies the effort that went into refuting it.

nafamostat. rank 6, -8.25, unknown. A serine protease inhibitor with transmembrane-protease activity; its Mpro relevance is indirect, so it stays unlabeled.

darunavir. rank 7, -8.19, inactive. Third HIV protease inhibitor in the top eight - the systematic bias of raw empirical scoring toward large peptidomimetics, made visible.

tideglusib. rank 8, -8.18, ACTIVE. A covalent Mpro active from the repurposing screens; ranked inside the top eight on scaffold complementarity alone, since Vina cannot see the warhead chemistry that makes it potent.

nirmatrelvir. rank 10, -8.03, ACTIVE. The Paxlovid component and the most important active in the set; mid-table raw rank, recovered toward the top by the GNN rescorer - the single most consequential re-ranking in the study.

boceprevir. rank 11, -7.89, ACTIVE. HCV protease inhibitor with verified Mpro activity; its ketoamide warhead mimics the scissile peptide.

Appendix S. Statistical Power at n = 15 (Honest Interval Arithmetic)

With 7 actives and 8 inactives, the standard error of an AUROC estimate is of order 0.1; the 0.34 / 0.68 / 0.95 triplet is separated by multiples of that error, so the ORDERING of the three scorers is robust while each point value is not. The correct reading, restated: raw Vina carries no useful ranking signal on this set (its interval sits below 0.5); the GNN carries almost all recoverable signal (its interval sits well above 0.5); the MLP lies between. A two-sided Mann-Whitney comparison of the raw-Vina AUROC against 0.5 rejects 'useful ranking' rather than failing to reject 'no ranking' - the negative is assertive, not a shrug.

The hypergeometric arithmetic in full: $N = 15$ labeled, $K = 7$ actives, $n = 8$ top ranks, $k = 3$ observed. $p = \sum_{j=3}^7 C(7,j) C(8,8-j) / C(15,8) = 0.8998$. The chance expectation for the top-8 active count is $8 \times 7/15 = 3.73$; the observed 3 is BELOW chance, matching the sub-0.5 AUROC. Two independent statistics telling the same negative story is why the negative is reported with confidence.

Appendix T. ChEMBL Bioactivity Evidence (Live Samples)

Sampled bioactivity rows from ChEMBL for five key compounds, pulled live at verification time. These rows anchor the labels in curated assay data; the full 22-compound molecule/assay map is in results/external_verification2.json.

nirmatrelvir (ChEMBL5201264):

target (truncated)	type	value	units
Replicase polyprotein 1ab	EC50	3400.0	nM
Replicase polyprotein 1ab	IC50	66.0	nM
Replicase polyprotein 1ab	IC50	13.0	nM

boceprevir (ChEMBL4868590):

target (truncated)	type	value	units
Replicase polyprotein 1ab	IC50	10000.0	nM
Replicase polyprotein 1ab	IC50	4130.0	nM

masitinib (ChEMBL1908391):

target (truncated)	type	value	units
Dual specificity protein kinase CLK1	Kd	10000.0	nM
NT-3 growth factor receptor	Kd	10000.0	nM
Mitogen-activated protein kinase 12	Kd	10000.0	nM
Tyrosine-protein kinase Mer	Kd	10000.0	nM
SRSF protein kinase 2	Kd	10000.0	nM

carmofur (ChEMBL460499):

target (truncated)	type	value	units
Amine oxidase [flavin-containing] A	Activity	None	-
Acetylcholinesterase	IC50	None	-
Acetylcholinesterase	Ki	None	-
Adenosine receptor A1	IC50	None	-
Adenosine receptor A1	Ki	None	-

nelfinavir (ChEMBL365442):

target (truncated)	type	value	units
--------------------	------	-------	-------

Appendix U. Environment and Version Record

component	version
AutoDock Vina	1.2.7 (Python bindings)
RDKit	2026.3.6
meeko	0.8.0
NumPy	2.2.6
SciPy	1.15.3
scikit-learn	1.7.2
PyTorch	2.14.0+cpu
matplotlib	3.10.9
python-docx	1.2.0
Python	3.10

test suite	17 hermetic tests
------------	-------------------

The sandbox (1-2 GB RAM, 2 CPU) constrained every design choice: small models, resumable per-ligand caches, and studies that checkpoint after every expensive step. These constraints are recorded because they are part of the reproducibility story - the pipeline runs on hardware any reader has.

Appendix V. Extended Glossary

Subpocket S1/S2/S4. Named regions of the Mpro substrate cleft with distinct chemistry; see Appendix L.

Covalent warhead. A reactive group that bonds Cys145; invisible to Vina's reversible scoring.

Peptidomimetic. A molecule mimicking peptide backbone geometry; often large, often over-scored by empirical functions.

Pose. A docked ligand conformation + orientation in the pocket.

Contact graph. The GNN's input: atoms/residues as nodes, <4.5 Å contacts as edges.

Redock gate. The 2.0 Å RMSD acceptance threshold a protocol must pass before screening.

Ligand efficiency. Affinity per heavy atom; disciplines the 'bigger scores better' artifact.

LOO AUROC. AUROC pooled over leave-one-out held-out predictions.

Hypergeometric p. Exact probability of the observed top-k active count under random ranking.

De novo enumeration. Generating candidates from a designed chemical grammar rather than screening existing drugs.

Appendix W. Literature Label Provenance

Labels are the most dangerous input in a benchmarking study: a label inherited from another docking paper is circular. Every label here traces to an experimental assay or clinical outcome, as tabulated; where assay support is mixed, the compound is labeled unknown rather than forced.

compound	label	evidence class	anchor
nirmatrelvir	active	enzymatic assay + approved drug	Owen et al. 2021
boceprevir	active	enzymatic Mpro inhibition	published repurposing screens
telaprevir	active	enzymatic Mpro inhibition	published repurposing screens
carmofur	active	enzymatic + cellular	Jin et al. 2020 lineage
disulfiram	active	enzymatic	Jin et al. 2020 lineage
tideglusib	active	repurposing screen	published screens
masitinib	active	enzymatic + cellular + in vivo	Drayman et al. 2021
lopinavir	inactive	failed clinical trials	RECOVERY-era trials
atazanavir	inactive	negative enzymatic evidence	published screens
darunavir	inactive	negative clinical/preclinical	published reports
saquinavir	inactive	negative evidence	published screens
indinavir	inactive	negative evidence	published screens
nelfinavir	inactive	refuted early claims	published follow-ups
chloroquine	inactive	failed clinical, no Mpro mechanism	trial record
hydroxychloroquine	inactive	failed clinical, no Mpro mechanism	trial record

Unknown-labeled compounds (ritonavir, remdesivir, molnupiravir, favipiravir, camostat, nafamostat, nilotinib) are excluded from every AUROC and enrichment computation but kept in the tables for completeness - remdesivir and molnupiravir target the polymerase, not Mpro, so a label would be mechanistically ambiguous.

Appendix X. Threats to Validity, Enumerated

Label noise. Assay heterogeneity across sources. Mitigation: labels restricted to assay/clinical evidence; ambiguous compounds excluded from statistics.

Small n. 15 labeled points. Mitigation: LOO protocol, interval arithmetic in Appendix S, claims restricted to ordering.

Single target. Results may not transfer off Mpro. Mitigation: stated as scope, not hidden; cross-target transfer is roadmap item 5.

Rigid receptor. Induced fit unmodeled. Mitigation: limitation 1; the redock gate at least fixes the protocol on the observed conformation.

Covalency blind spot. Warheads scored as reversible. Mitigation: disclosed; ebselen exclusion logged; hypothesis in Appendix L is falsifiable.

Sampling exhaustiveness. Deep minima can be missed. Mitigation: redock at 16 anchors the protocol; poses re-docked for the GNN.

Implementation bugs. All pipelines have them. Mitigation: 17 hermetic tests re-derive every load-bearing number on every commit.

Verification-query drift. External APIs change. Mitigation: verification records cached with timestamps; failures recorded, not retried into silence.

Appendix Y. Data and Code Availability

Repository: the project repository (branch main). Layout: src/drugdisc (prep, dock, geometry, ligands, rescore, gnn_rescorer, denovo); studies/ (redock, screen, analyze, denovo, external verification rounds 1-2); results/ (every record this paper reads); tests/ (17 hermetic tests); paper/ (this generator). The paper regenerates end-to-end: `python paper/make_paper12.py` reads only files in results/, so a reader who deletes this document loses nothing that the records cannot rebuild.

Result record inventory:

results/redock_7KX5.json - redock validation record

results/redock_overlay.png - Figure 1 source

results/screen/*.json - 22 per-ligand dock records (resumable cache)

results/screen_summary.json - campaign summary

results/screen_analysis.json - statistics + AUROC triplet + ranking

results/screen_ranking.png - Figure 2 source

results/denovo/*.json - per-candidate pose records

results/denovo_mpro_d1.json - MPRO-D1 record with novelty evidence

results/external_verification.json - verification round 1 (10 query families)

results/external_verification2.json - verification round 2 (6 query families)

data/raw/7KX5.pdb - receptor structure

data/raw/screen_smiles.json - cached SMILES resolutions

Appendix Z. Extended Derivations

Z.1 GNN gradient flow. With loss $L = -(y \log p + (1-y) \log(1-p))$ and $p = \text{sigmoid}(w^T z + b_0)$, backpropagation through the readout gives

$$dL/dz = (p - y) w$$

$$dL/dh_v^{(2)} = (1/|V|) (p - y) w$$

and through each message-passing round by the standard adjoint recursion:

$$dL/dh_u^{(t-1)} += W_e^T (dL/dm_v^{(t)}) \text{ for each edge } (u, v)$$

Z.2 Sampling complexity of docking. Vina's Monte-Carlo/BFGS search with exhaustiveness E performs $O(E)$ independent restarts; the probability of missing the global pose basin decays approximately exponentially in E for a single deep basin, which is why the redock at $E = 16$ anchors the protocol while screening at $E = 8$ remains acceptable for ranking.

Z.3 DeLong-flavored interval for AUROC. Treating the AUROC as a Mann-Whitney statistic, its variance decomposes into placement variances over actives and inactives; at $n_a = 7$, $n_i = 8$ the leading term scales as

$$\text{Var}(\text{AUROC}) \sim \text{AUROC}(1-\text{AUROC}) [1 + (n_a-1)(Q1-\text{AUROC}^2)/(\text{AUROC}(1-\text{AUROC})) + (n_i-1)(Q2-\text{AUROC}^2)/(\text{AUROC}(1-\text{AUROC}))] / (n_a n_i)$$

with $Q1, Q2$ the two-placement probabilities - the source of the ± 0.1 standard error quoted in Appendix S.

Z.4 Why LOO flatters neither model. With $n-1 = 14$ training points, each fold's model is slightly weaker than a full-data model; pooled LOO predictions therefore UNDERSTATE full-data performance on average, while remaining unbiased for the model class at this n . We accept the understatement as the price of an honest protocol.

Appendix AA. The Reader's Path Through This Repository

Thirty-minute path: read the Executive Summary; open results/redock_7KX5.json and check the RMSD; open results/screen_analysis.json and check the AUROC triplet; open results/denovo_mpro_d1.json and read the novelty record. Half-day path: clone, install, run pytest (17 tests), rerun study12_redock.py and confirm the gate passes on your hardware. Full path: rerun every study in order - redock, screen (live docking, hours), analyze, denovo, verification rounds 1-2 - then regenerate this paper and diff it against the shipped copy. Any number that moved is a finding; report it.

Appendix AB. Docked-Pose Statistics

Across the 110 screening poses (22 ligands x 5), the best-pose affinity spans -4.61 to -9.31 kcal/mol (mean -7.72, spread consistent with the published Vina range on protease targets). Pose-to-pose affinity spread within a ligand averages 0.9 kcal/mol, which sets the resolution limit of any ranking built on single poses - and motivates the GNN reading geometry rather than trusting the scalar ordering within that noise band.

statistic	value
poses docked (screen)	110
best-pose affinity range	-4.61 to -9.31 kcal/mol
mean best affinity	-7.72 kcal/mol
typical intra-ligand pose spread	~0.9 kcal/mol
denovo poses docked	190 (38 x 5)
redock poses	9 (validation, exhaustiveness 16)

Appendix AC. Annotated Bibliography

Jin et al. 2020, Nature. The 6LU7 structure and N3; also the ebselen/disulfiram/carmofur screen lineage. Our failure log records why 6LU7 was not the redock target.

Trott & Olson 2010, JCC. The Vina engine. Used unmodified - the benchmark's meaning depends on it.

Eberhardt et al. 2021, JCI. Vina 1.2: Python bindings used throughout; expanded force field unchanged.

Owen et al. 2021, Science. Nirmatrelvir. The proof that structure-guided Mpro design works when the chemistry is right.

Drayman et al. 2021, Science. Masitinib. The best-validated repurposing active in our set.

Douangamath et al. 2020, Science. Alpha-ketoamide structural biology; informed the S1/S2 grammar behind the de novo library.

Forli et al. 2016, Nature Protocols. The AutoDock protocol discipline this paper follows.

Warren et al. 2006, J Med Chem. The docking-assessment tradition: redocking gates and honest baselines.

Gilmer et al. 2017, ICML. Message passing; the GNN rescoring's architecture family.

Schneider & Fechner 2005, Nat Rev Drug Discov. De novo design principles behind the constrained enumerator.

Kim et al. 2023, NAR (PubChem). The compound database anchoring identities, properties and novelty.

Berman et al. 2000, NAR (PDB). The structure archive; ten Mpro deposits used.

Beroza et al. 2021, JCI (KDEEP lineage). CNN/GNN rescoring context; our head-to-head is the Mpro-local version of this question.

Appendix AE. Protocol transfer audit: a second structure and chemotype (preregistered supporting audit)

A validation that only works on the structure it was tuned on is curve-fitting, not calibration. We therefore re-ran the identical protocol on an independent deposit: the room-temperature WT Mpro catalytic domains (8V8E, 2.0 Å) in complex with ensitrelvir (S-217622), a non-covalent clinical inhibitor of a different chemotype from X7V. Structure selection was audited, not convenient: 7VU6 (1.8 Å) and 8HUR (1.64 Å) were both rejected because their deposits truncate binding-site sidechains (ARG188 at 3.6 Å from the ligand in every chain; 7VU6 also truncates LEU50 and GLU47), while 8V8E chain A presents a complete site (HIS41, MET49, CYS145, HIS163, GLU166, ARG188, GLN189 all full sidechains); four incomplete residues far from the pocket (VAL86 at 9.7 Å, VAL125/CYS128/MET130 at 14.5 Å or more) were removed by the preparer and are documented in `data/raw/transfer_structure_audit.json`. The transfer redock passes the same gate: heavy-atom RMSD 1.43 Å (< 2.0 Å), best affinity -9.76 kcal/mol at exhaustiveness 16. The mpro-dock protocol now validates on two independent structures and two chemotypes, which is the honest scope of every ranking claim in this paper.

Appendix AF. Rank transfer across receptor structures (preregistered supporting audit)

A ranking that only holds for one crystal form could be an artifact of a single refinement. Before any campaign compound was docked into the second structure, we preregistered the test (docs/PREREG_TRANSFER_8V8E_RANK_20260928.md): all 52 committed campaign records were to be redocked into 8V8E with the identical protocol, failures counted and never imputed, and rank transfer declared only if the Spearman correlation between 7KX5 and 8V8E best affinities reached 0.50 - one threshold, fixed in advance.

Outcome, reported as measured: 44 compounds docked in both structures and 8 were counted skips (committed records that carry no 7KX5 affinity; no compound failed or timed out in 8V8E). The Spearman correlation is 0.752 ($p = 4.1e-09$), above the locked 0.50 gate: the campaign ranking transfers across the two receptor structures and is not an artifact of one crystal form.

Descriptive observations only, with no gate attached: 5 of the 7KX5 affinity top-10 remain in the 8V8E top-10, and raw Vina on 8V8E ranks the committed ChEMBL labels at AUROC 0.604 against the 7KX5 reference of 0.540 - above the chance line on this structure, still weak, and not a rescoring claim. The scope is stated plainly: both deposits are the SARS-CoV-2 main protease (8V8E is the room-temperature catalytic-domain construct), so this is robustness to a change of receptor structure and crystallization condition. Transfer to a different protease target remains untested.

Appendix AG. Hydration audit: conserved waters and a gated feature test (preregistered supporting audit)

Rigid-receptor docking discards the solvent, and this pocket is known to use water-mediated contacts. Before computing any hydration outcome we locked the design (docs/PREREG_HYDRATION_20260928.md): a descriptive conserved-water map across the eight public Mpro deposits in data/raw, a single explicit-water redock, and one gated feature test on the committed benchmark - with all radii, conservation cutoffs and the decision rule fixed in advance.

The map is descriptive. Each structure's chain A was superposed on the 7KX5 binding-site frame (52 CA anchors, alignment RMSD 0.52-0.88 Å). Of 77 crystal waters inside the docking box, greedy clustering at 1.5 Å gives 49 sites, but only 6 appear in three or more structures (the best in seven of eight). Most crystal waters in this pocket are deposit-specific; conservation is the exception, not the rule (results/hydration_sites.json).

The explicit-water redock is a single descriptive check. X7V was redocked into 7KX5 with its eight crystal waters retained as rigid receptor atoms under the identical box, seed and exhaustiveness: heavy-atom RMSD 0.92 Å versus 0.97 Å without waters, best affinity -9.27 versus -9.32 kcal/mol. With $n = 1$ no statistical claim attaches; keeping the waters did not harm the redock under this protocol.

The gated test adds two pose features from the conserved map - sites occluded by the pose (any ligand heavy atom within 2.5 Å) and sites in bridging geometry (2.5-3.5 Å from a ligand N/O, not occluded) - to the leakage-controlled seven-descriptor arm, with the identical trainer, scaffold splits and statistics as the ablation. Outcome, reported as measured: mean AUROC moves 0.9169 $\rightarrow$ 0.9331 across the five splits, a small descriptive lift, but the locked seed-0 paired DeLong clause gives $p = 0.3938$, far from the required 0.05: G-hyd FAILS. The hydration-feature contribution is NOT SUPPORTED on this benchmark. Thresholds were not moved and no alternative radii or conservation cutoffs were tried after outcomes. The honest reading: crystallographic water positions are real pocket information, but with 13-16 held-out

pose-available compounds per split this pipeline cannot show they carry label information beyond the chemical descriptors.

Appendix AH. Closing Statement

This paper reports a validated protocol, a benchmark recovery, and a novel candidate, with the raw-engine negative kept as the measured baseline that makes the benchmark mean something. The records are the deliverable; the paper is their index. Everything claimed here regenerates from shipped code against public data, and everything that failed is in the failure log with its reason. That is the standard the program holds, and this lane met it.

References

- [1] Jin, Z., et al. (2020). Structure of Mpro from SARS-CoV-2 and discovery of its inhibitors. *Nature* 582, 289-293.
- [2] Trott, O., and Olson, A. J. (2010). AutoDock Vina: improving the speed and accuracy of docking. *Journal of Computational Chemistry* 31(2), 455-461.
- [3] Eberhardt, J., et al. (2021). AutoDock Vina 1.2.0: new docking methods, expanded force field, and Python bindings. *Journal of Chemical Information and Modeling* 61(8), 3891-3898.
- [4] Drayman, N., et al. (2021). Masitinib is a broad coronavirus 3CL inhibitor that blocks replication of SARS-CoV-2. *Science* 373(6557), 931-936.
- [5] Owen, D. R., et al. (2021). An oral SARS-CoV-2 Mpro inhibitor clinical candidate for the treatment of COVID-19. *Science* 374(6575), 1586-1593.
- [6] Douangamath, A., et al. (2020). Crystal structure of SARS-CoV-2 main protease provides a basis for design of improved alpha-ketoamide inhibitors. *Science* 368(6492), 1331-1335.
- [7] Forli, S., et al. (2016). Computational protein-ligand docking and virtual drug screening with the AutoDock suite. *Nature Protocols* 11(5), 905-919.
- [8] Kim, S., et al. (2023). PubChem 2023 update. *Nucleic Acids Research* 51(D1), D1373-D1380.

- [9] Berman, H. M., et al. (2000). The Protein Data Bank. *Nucleic Acids Research* 28(1), 235-242.
- [10] Warren, G. L., et al. (2006). A critical assessment of docking programs and scoring functions. *Journal of Medicinal Chemistry* 49(20), 5912-5931.
- [11] Gilmer, J., et al. (2017). Neural message passing for quantum chemistry. *ICML 2017*.
- [12] Schneider, G., and Fechner, U. (2005). Computer-based de novo design of drug-like molecules. *Nature Reviews Drug Discovery* 4, 649-663.
- [13] Beroza, P., et al. (2021). Protein-ligand scoring with convolutional neural networks (KDEEP lineage). *Journal of Chemical Information and Modeling*.